



Concept Learning




Waiting outside the house to get an autograph.



Which days does he come out to enjoy sports?

- Sky condition
 - Humidity
 - Temperature
 - Wind
 - Water
 - Forecast
-
- Attributes of a day: takes on values



Learning Task

- We want to make a hypothesis about the day on which SRK comes out..
 - in the form of a boolean function on the attributes of the day.
- Find the right hypothesis/function from historical data

Training Examples for EnjoySport

	Sky	Temp	Humid	Wind	Water	Forecst	EnjoySpt
c(Sunny	Warm	Normal	Strong	Warm	Same)=1	Yes
c(Sunny	Warm	High	Strong	Warm	Same)=1	Yes
c(Rainy	Cold	High	Strong	Warm	Change)=0	No
c(Sunny	Warm	High	Strong	Cool	Change)=1	Yes

■ Negative and positive learning examples

■ Concept learning: c is the target concept

- Deriving a Boolean function from training examples

- Many “hypothetical” boolean functions

➤ Hypotheses; find h such that $h = c$.

– Other more complex examples:

❖ Non-boolean functions

■ Generate hypotheses for concept from TE's

Representing Hypotheses

- Task of finding appropriate set of hypotheses for concept given training data
- Represent hypothesis as **Conjunction** of **constraints** of the following form:
 - Values possible in any hypothesis
 - ♦ Specific value : Water = *Warm*
 - ♦ Don't-care value: Water = ?
 - ♦ No value allowed : Water = \emptyset
 - i.e., no permissible value given values of other attributes
 - Use vector of such values as hypothesis:
 - ♦ $\langle \text{Sky} \quad \text{AirTemp} \quad \text{Humid} \quad \text{Wind} \quad \text{Water} \quad \text{Forecast} \rangle$
 - Example: $\langle \text{Sunny} \quad ? \quad ? \quad \text{Strong} \quad ? \quad \text{Same} \rangle$
- Idea of *satisfaction of hypothesis* by some example
 - say “example satisfies hypothesis”
 - defined by a function $h(x)$:
$$h(x) = 1 \text{ if } h \text{ is true on } x$$
$$= 0 \text{ otherwise}$$
- Want hypothesis that best fits examples:
 - Can reduce learning to search problem over space of hypotheses

Prototypical Concept Learning Task

TASK T: predicting when person will enjoy sport

- **Target function** c : $\text{EnjoySport} : X \rightarrow \{0, 1\}$
- Cannot, in general, know Target function c
 - ❖ Adopt hypotheses H about c
- Form of hypotheses H :
 - ❖ Conjunctions of literals $\langle ?, \text{Cold}, \text{High}, ?, ?, ? \rangle$

■ EXPERIENCE E

- **Instances** X : possible days described by attributes *Sky, AirTemp, Humidity, Wind, Water, Forecast*
- **Training examples** D : Positive/negative examples of target function $\{\langle x_1, c(x_1) \rangle, \dots, \langle x_m, c(x_m) \rangle\}$

■ **PERFORMANCE MEASURE P:** Hypotheses h in H such that $h(x) = c(x)$ for all x in D ()

- There may exist several alternative hypotheses that fit examples

Inductive Learning Hypothesis

Any hypothesis found to approximate the target function well over a sufficiently large set of training examples will also approximate the target function well over other unobserved examples

Approaches to learning algorithms

- Brute force search

The choice of the hypothesis space reduces the number of hypotheses.

- Enumerate all possible hypotheses and evaluate
- Highly inefficient even for small *EnjoySport* example
 - ♦ $|X| = 3.2.2.2.2 = 96$ distinct *instances*
 - ♦ Large number of *syntactically distinct* hypotheses (0's, ?'s)
 - EnjoySport: $|H| = 5.4.4.4.4 = 5120$
 - Fewer when consider h's with 0's

Every h with a 0 is empty set of instances (classifies instance as neg)

Hence # semantically distinct h's is:

$$1 + (4.3.3.3.3) = 973$$


- ♦ EnjoySport is VERY small problem compared to many

- Hence use other search procedures.

- Approach 1: Search based on ordering of hypotheses
- Approach 2: Search based on finding all possible hypotheses using a good representation of hypothesis space
 - ♦ All hypotheses that fit data

Ordering on Hypotheses


Instances X



$x_1 = \langle \text{Sunny Warm High Strong Cool Same} \rangle$

$x_2 = \langle \text{Sunny Warm High Light Warm Same} \rangle$

Hypotheses H



↑ specific
↓ general

$h_1 = \langle \text{Sunny ? ? Strong ? ?} \rangle$

$h_2 = \langle \text{Sunny ? ? ? ? ?} \rangle$

$h_3 = \langle \text{Sunny ? ? ? Cool ?} \rangle$

- h is more general than h' ($h \geq_g h'$) if for each instance x ,

$$h'(x) = 1 \rightarrow h(x) = 1$$

- Which is the most general/most specific hypothesis?

Find-S Algorithm

Assumes

- There is hypothesis h in H describing target function c
- There are no errors in the TEs

Procedure

1. Initialize h to the most specific hypothesis in H (*what is this?*)
2. For each *positive* training instance x
 - For each attribute constraint a_i in h
 - If the constraint a_i in h is satisfied by x
 - do nothing
 - Else
 - replace a_i in h by the next more general constraint that is satisfied by x
3. Output hypothesis h

Note


There is no change for a negative example, so they are ignored.

This follows from assumptions that there is h in H describing target function c (*ie., for this h , $h=c$*) *and* that there are no errors in data. In particular, it follows that the hypothesis at any stage cannot be changed by neg example.


Assumption: Everything except the positive examples is negative

Example of Find-S

Instances X



Hypotheses H



↑ specific
↓ general

$x_1 = \langle \text{Sunny Warm Normal Strong Warm Same} \rangle +$
 $x_2 = \langle \text{Sunny Warm High Strong Warm Same} \rangle +$
 $x_3 = \langle \text{Rainy Cold High Strong Warm Change} \rangle -$
 $x_4 = \langle \text{Sunny Warm High Strong Cool Change} \rangle +$

$h_0 = \langle \emptyset \ \emptyset \ \emptyset \ \emptyset \ \emptyset \ \emptyset \rangle$

$h_1 = \langle \text{Sunny Warm Normal Strong Warm Same} \rangle$

$h_2 = \langle \text{Sunny Warm ? Strong Warm Same} \rangle$

$h_3 = \langle \text{Sunny Warm ? Strong Warm Same} \rangle$

$h_4 = \langle \text{Sunny Warm ? Strong ? ?} \rangle$

Problems with Find-S

- Problems:
 - Throws away information!
 - ♦ Negative examples
 - Can't tell whether it has learned the concept
 - ♦ Depending on H , there might be several h 's that fit TEs!
 - ♦ Picks a maximally specific h (why?)
 - Can't tell when training data is inconsistent
 - ♦ Since ignores negative TEs
- But
 - It is simple
 - Outcome is independent of order of examples
 - ♦ Why?
- What alternative overcomes these problems?
 - Keep *all* consistent hypotheses!
 - ♦ Candidate elimination algorithm

Consistent Hypotheses and Version Space

- A hypothesis h is **consistent** with a set of training examples D of target concept c if $h(x) = c(x)$ for each training example $\langle x, c(x) \rangle$ in D
 - Note that consistency is with respect to specific D .

- Notation:

$$\text{Consistent}(h, D) \equiv \forall \langle x, c(x) \rangle \in D :: h(x) = c(x)$$

- The **version space**, $VS_{H,D}$, with respect to hypothesis space H and training examples D , is the subset of hypotheses from H consistent with D

- Notation:


$$VS_{H,D} = \{ h \mid h \in H \wedge \text{Consistent}(h, D) \}$$

List-Then-Eliminate Algorithm


1. $VersionSpace \leftarrow$ list of all hypotheses in H
2. For each training example $\langle x, c(x) \rangle$
remove from $VersionSpace$ any hypothesis h for which
 $h(x) \neq c(x)$
3. Output the list of hypotheses in $VersionSpace$
4. This is essentially a brute force procedure

Example of Find-S, Revisited

Instances X



Hypotheses H



specific
↑
↓
general

$x_1 = \langle \text{Sunny Warm Normal Strong Warm Same} \rangle +$
 $x_2 = \langle \text{Sunny Warm High Strong Warm Same} \rangle +$
 $x_3 = \langle \text{Rainy Cold High Strong Warm Change} \rangle -$
 $x_3 = \langle \text{Sunny Warm High Strong Cool Change} \rangle +$

$h_0 = \langle \emptyset \emptyset \emptyset \emptyset \emptyset \emptyset \rangle$


$h_1 = \langle \text{Sunny Warm Normal Strong Warm Same} \rangle$

$h_2 = \langle \text{Sunny Warm ? Strong Warm Same} \rangle$

$h_3 = \langle \text{Sunny Warm ? Strong Warm Same} \rangle$

$h_4 = \langle \text{Sunny Warm ? Strong ? ?} \rangle$

Version Space for this Example



Representing Version Spaces

- Want more compact representation of VS
 - Store most/least general boundaries of space
 - Generate all intermediate h's in VS
 - Idea that any h in VS must be consistent with all TE's
 - ♦ Generalize from most specific boundaries
 - ♦ Specialize from most general boundaries
- The **general boundary**, G , of version space $VS_{H,D}$ is the set of its maximally general members consistent with D
 - Summarizes the negative examples; anything more general will cover a negative TE
- The **specific boundary**, S , of version space $VS_{H,D}$ is the set of its maximally specific members consistent with D
 - Summarizes the positive examples; anything more specific will fail to cover a positive TE

Theorem

Every member of the version space lies between the S,G boundary

$$VS_{H,D} = \{h \mid h \in H \wedge \exists s \in S \exists g \in G (g \geq h \geq s)\}$$

- Must prove:
 - 1) every h satisfying RHS is in $VS_{H,D}$;
 - 2) every member of $VS_{H,D}$ satisfies RHS.
- For 1), let g, h, s be arbitrary members of G, H, S respectively with $g > h > s$
 - s must be satisfied by all + TEs and so must h because it is more general;
 - g cannot be satisfied by any – TEs, and so nor can h
 - h is in $VS_{H,D}$ since satisfied by all + TEs and no – TEs
- For 2),
 - Since h satisfies all + TEs and no – TEs, $h \geq s$, and $g \geq h$.

Candidate Elimination Algorithm

$G \leftarrow$ maximally general hypotheses in H

$S \leftarrow$ maximally specific hypotheses in H

For each training example d , do

- If d is positive
 - Remove from G every hypothesis inconsistent with d
 - For each hypothesis s in S that is inconsistent with d
 - ♦ Remove s from S
 - ♦ Add to S all minimal generalizations h of s such that
 1. h is consistent with d , and
 2. some member of G is more general than h
 - Remove from S every hypothesis that is more general than another hypothesis in S

Candidate Elimination Algorithm (cont)

- If d is a negative example
 - Remove from S every hypothesis inconsistent with d
 - For each hypothesis g in G that is inconsistent with d
 - ♦ Remove g from G
 - ♦ Add to G all minimal specializations h of g such that
 1. h is consistent with d , and
 2. some member of S is more specific than h
 - Remove from G every hypothesis that is less general than another hypothesis in G
- Essentially use
 - Pos TEs to generalize S
 - Neg TEs to specialize G
- Independent of order of TEs
- Convergence guaranteed if:
 - ***no errors***
 - ***there is h in H describing c .***

Example

Recall : If d is positive

Remove from G every hypothesis inconsistent with d

For each hypothesis s in S that is inconsistent with d

- Remove s from S
- Add to S all minimal generalizations h of s that are specializations of a hypothesis in G
- Remove from S every hypothesis that is more general than another hypothesis in S

S_0 $\{\langle \emptyset \ \emptyset \ \emptyset \ \emptyset \ \emptyset \ \emptyset \rangle\}$

G_0 $\{\langle ? \ ? \ ? \ ? \ ? \ ? \rangle\}$

$\langle \text{Sunny Warm Normal Strong Warm Same} \rangle$ +

S_1 $\{\langle \text{Sunny Warm Normal Strong Warm Same} \rangle\}$

G_1 $\{\langle ? \ ? \ ? \ ? \ ? \ ? \rangle\}$

Example (contd)

S_1 {⟨Sunny Warm Normal Strong Warm Same⟩}

G_1 {⟨? ? ? ? ? ?⟩}

⟨Sunny Warm High Strong Warm Same⟩ +

S_2 {⟨Sunny Warm ? Strong Warm Same⟩}

G_2 {⟨? ? ? ? ? ?⟩}

Example (contd)



S_2 {⟨Sunny Warm ? Strong Warm Same⟩}

Recall: If d is a negative example

G_2 {⟨? ? ? ? ? ?⟩}

- Remove from S every hypothesis inconsistent with d
- For each hypothesis g in G that is inconsistent with d
 - ❖ Remove g from G
 - ❖ Add to G all minimal specializations h of g that generalize some hypothesis in S
 - ❖ Remove from G every hypothesis that is less general than another hypothesis in G

⟨Rainy Cold High Strong Warm Change⟩ -

S_3 {⟨Sunny Warm ? Strong Warm Same⟩}

Current G boundary is incorrect
So, need to make it more specific.

G_3 {⟨Sunny ? ? ? ? ?⟩, ⟨? Warm ? ? ? ?⟩, ⟨? ? ? ? ? Same⟩}

Example (contd)

- Why are there no hypotheses left relating to:
□ $\langle \text{Cloudy ? ? ? ? ?} \rangle$
- The following specialization using the third value
 $\langle ? ? \text{Normal ? ? ?} \rangle$,
is not more general than the specific boundary

$\{ \langle \text{Sunny Warm ? Strong Warm Same} \rangle \}$
- The specializations $\langle ? ? ? \text{Weak ? ?} \rangle$,
 $\langle ? ? ? ? \text{Cool ?} \rangle$ are also inconsistent with S

Example (contd)

$$S_3 \quad \{ \langle \text{Sunny Warm ? Strong Warm Same} \rangle \}$$

$$G_3 \quad \{ \langle \text{Sunny ? ? ? ? ?} \rangle, \langle \text{? Warm ? ? ? ?} \rangle, \langle \text{? ? ? ? ? Same} \rangle \}$$

$\langle \text{Sunny Warm High Strong Cool Change} \rangle$ +

$$S_4 \quad \{ \langle \text{Sunny Warm ? Strong ? ?} \rangle \}$$

$$G_4 \quad \{ \langle \text{Sunny ? ? ? ? ?} \rangle, \langle \text{? Warm ? ? ? ?} \rangle \}$$

Example (contd)

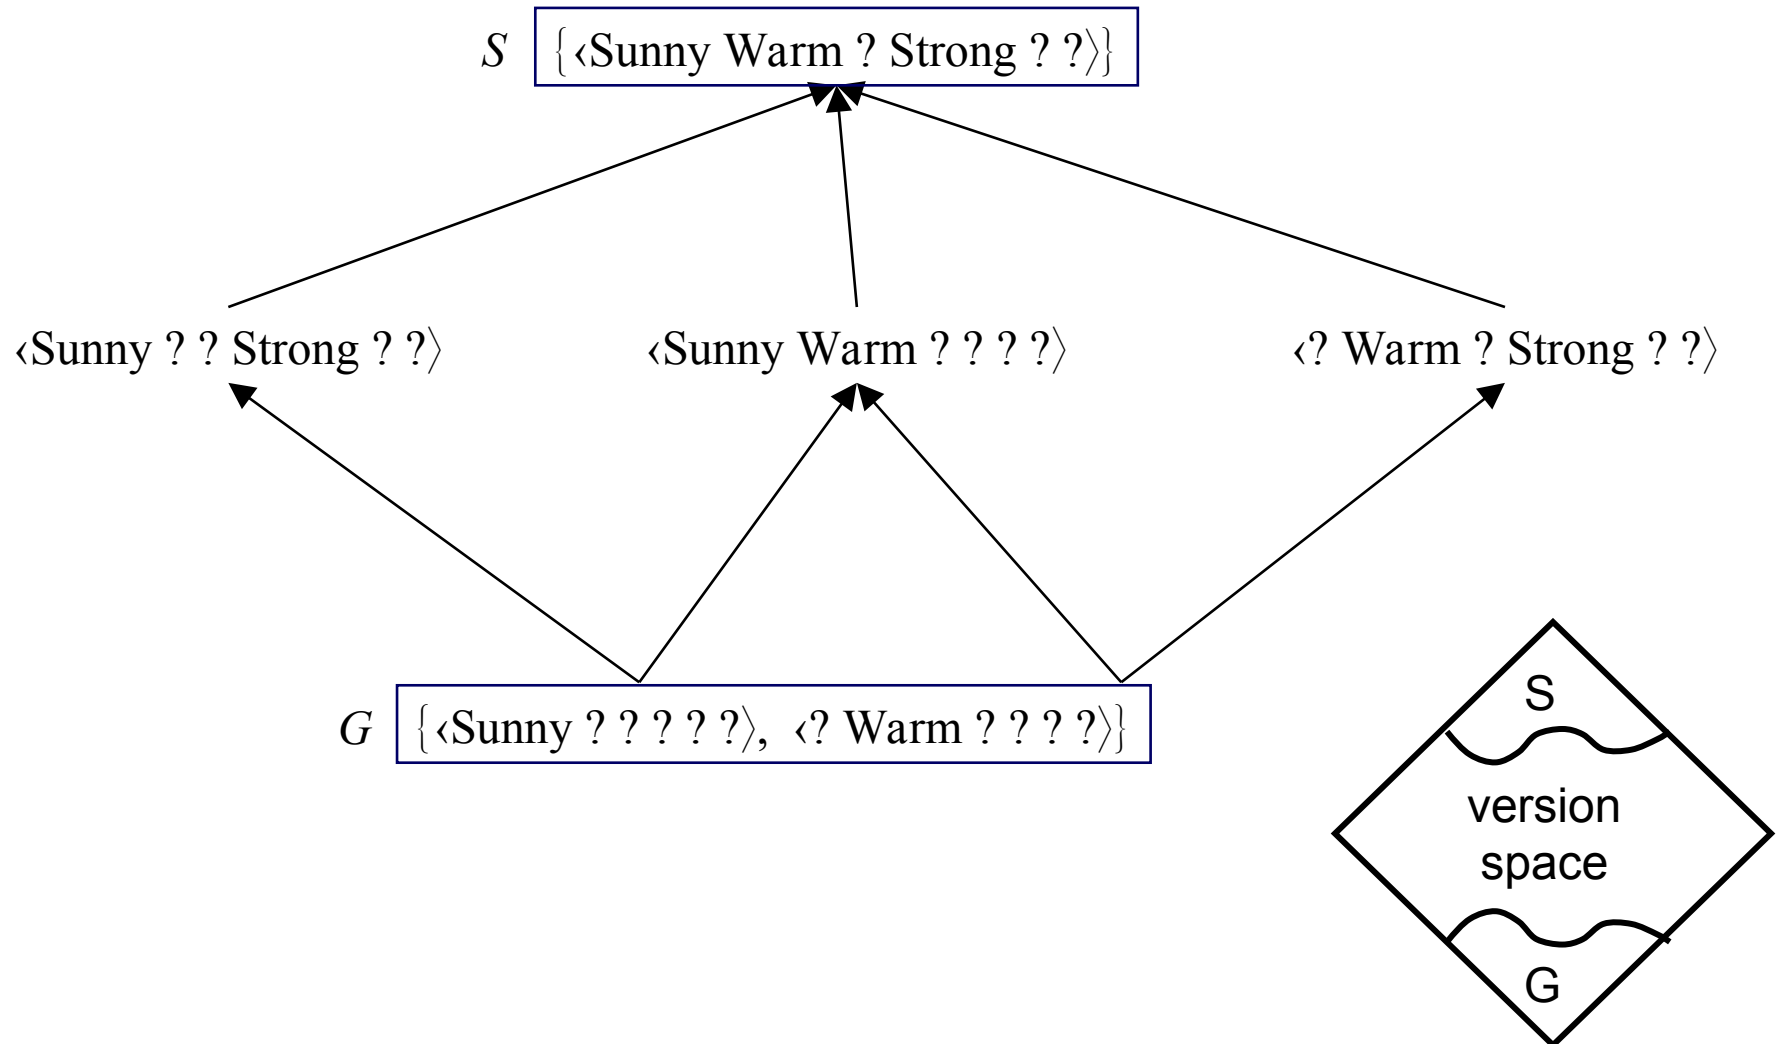
⟨*Sunny Warm High Strong Cool Change*⟩ +

- Why does this example remove a hypothesis from G?:

□ ⟨? ? ? ? ? *Same*⟩

- This hypothesis
 - Cannot be specialized, since would not cover new TE
 - Cannot be generalized, because more general would cover negative TE.
 - Hence must drop hypothesis.

Version Space of the Example



Convergence of algorithm

- Convergence guaranteed if:
 - *no errors*
 - *there is h in H describing c .*
- Ambiguity removed from VS when $S = G$
 - Containing single h
 - When have seen enough TEs
- If have false negative TE, algorithm will remove every h consistent with TE, and hence will remove correct target concept from VS
 - If observe enough TEs will find that S, G boundaries converge to empty VS


Let us try this

Origin	Manufacturer	Color	Decade	Type	
Japan	Honda	Blue	1980	Economy	+
Japan	Toyota	Green	1970	Sports	-
Japan	Toyota	Blue	1990	Economy	+
USA	Chrysler	Red	1980	Economy	-
Japan	Honda	White	1980	Economy	+

And this

Origin	Manufacturer	Color	Decade	Type	
Japan	Honda	Blue	1980	Economy	+
Japan	Toyota	Green	1970	Sports	-
Japan	Toyota	Blue	1990	Economy	+
USA	Chrysler	Red	1980	Economy	-
Japan	Honda	White	1980	Economy	+
Japan	Toyota	Green	1980	Economy	+
Japan	Honda	Red	1990	Economy	-

Which Next Training Example?




Assume learner can choose the next TE

- Should choose d such that
 - Reduces maximally the number of hypotheses in VS
 - Best TE: satisfies precisely 50% hypotheses;
 - ♦ Can't always be done
 - Example:
 - ♦ $\langle \text{Sunny Warm Normal Weak Warm Same} \rangle ?$
 - ♦ If pos, generalizes S
 - ♦ If neg, specializes G

Order of
examples matters
for intermediate
sizes of S, G ; not
for the final S, G

Classifying new cases using VS



- Use *voting procedure* on following examples:

- ⟨Sunny Warm Normal Strong Cool Change⟩
- ⟨Rainy Cool Normal Weak Warm Same⟩
- ⟨Sunny Warm Normal Weak Warm Same⟩
- ⟨Sunny Cold Normal Strong Warm Same⟩

Effect of incomplete hypothesis space

- Preceding algorithms work if target function is in H
 - Will generally not work if target function *not* in H
- Consider following examples which represent target function

“sky = sunny or sky = cloudy”:

 - ⟨Sunny Warm Normal Strong Cool Change⟩ Y
 - ⟨Cloudy Warm Normal Strong Cool Change⟩ Y
 - ⟨Rainy Warm Normal Strong Cool Change⟩ N
- If apply CE algorithm as before, end up with empty VS
 - After first two TEs, S= ⟨? Warm Normal Strong Cool Change⟩
 - New hypothesis is overly general
 - ♦ it covers the third negative TE!
- Our H does not include the appropriate c

Need more
expressive
hypotheses

Incomplete hypothesis space

- If c not in H , then consider generalizing representation of H to contain c
 - For example, add disjunctions or negations to representation of hypotheses in H
- One way to avoid problem is to allow **all** possible representations of h 's
 - Equivalent to allowing all possible subsets of instances as defining the concept of EnjoySport
 - ♦ Recall that there are 96 instances in EnjoySport; hence there are 2^{96} possible hypotheses in full space H
 - ♦ Can do this by using full propositional calculus with AND, OR, NOT
 - ♦ Hence H defined only by conjunctions of attributes is biased (containing only 973 h 's)

Unbiased Learners and Inductive Bias

- BUT if have no limits on representation of hypotheses (i.e., full logical representation: *and*, *or*, *not*), can only learn examples...no generalization possible!
 - Say have 5 TEs $\{x_1, x_2, x_3, x_4, x_5\}$, with x_4, x_5 negative TEs
- Apply CE algorithm
 - S will be disjunction of positive examples ($S = \{x_1 \text{ OR } x_2 \text{ OR } x_3\}$)
 - G will be negation of disjunction of negative examples ($G = \{\text{not } (x_4 \text{ or } x_5)\}$)
 - Need to use all instances to learn the concept!
- Cannot predict usefully:
 - TEs have unanimous vote
 - other h 's have 50/50 vote!
 - ♦ For every h in H that predicts +, there is another that predicts -

Unbiased Learners and Inductive Bias

- Approach:
 - Place constraints on representation of hypotheses
 - ♦ Example of limiting connectives to conjunctions
 - ♦ Allows learning of generalized hypotheses
 - ♦ Introduces bias that depends on hypothesis representation
- Need formal definition of inductive bias of learning algorithm

Inductive Syst and Equiv Deductive Syst

- Inductive bias made explicit in *equivalent deductive system*
 - Logically represented system that produces same outputs (classification) from inputs (TEs, instance x , bias B) as CE procedure
- Inductive bias (IB) of learning algorithm L is any minimal set of assertions B such that for any target concept c and training examples D , we can logically infer value $c(x)$ of any instance x from B , D , and x
 - E.g., for rote learner, $B = \{\}$, and there is no IB
- Difficult to apply in many cases, but a useful guide

Inductive Bias and specific learning algs

- Rote learners:
no IB
- Version space candidate elimination algorithm:
c can be represented in H
- Find-S: c can be represented in H;
all instances that are not positive are negative

Computational Complexity of VS

- The S set for conjunctive feature vectors and tree-structured attributes is linear in the number of features and the number of training examples.
- The G set for conjunctive feature vectors and tree-structured attributes can be exponential in the number of training examples.
- In more expressive languages, both S and G can grow exponentially.
- The order in which examples are processed can significantly affect computational complexity.

Exponential size of G

- n Boolean attributes
- 1 positive example: (T, T, \dots, T)
- $n/2$ negative examples:
 - (F, F, T, \dots, T)
 - $(T, T, F, F, T, \dots, T)$
 - $(T, T, T, T, F, F, T, \dots, T)$
 - ..
 - (T, \dots, T, F, F)
- Every hypothesis in G needs to choose from $n/2$ 2-element sets.
 - Number of hypotheses $= 2^{n/2}$

Summary

- Concept learning as search through H
- General-to-specific ordering over H
- Version space candidate elimination algorithm
- S and G boundaries characterize learner's uncertainty
- Learner can generate useful queries
- Inductive leaps possible only if learner is biased!
- Inductive learners can be modeled as equiv deductive systems
- Biggest problem is inability to handle data with errors
 - Overcome with procedures for learning decision trees